

Untitled2

December 27, 2019

```
[2]: import numpy as np
import pandas as pd
```

```
[54]: # Import clean data
path = 'kc_house_data.csv'
df_H = pd.read_csv(path)
df_H.head()
```

```
[54]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
0	7129300520	20141013T000000	221900.0	3	1.00	1180	
1	6414100192	20141209T000000	538000.0	3	2.25	2570	
2	5631500400	20150225T000000	180000.0	2	1.00	770	
3	2487200875	20141209T000000	604000.0	4	3.00	1960	
4	1954400510	20150218T000000	510000.0	3	2.00	1680	

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	\
0	5650	1.0	0	0	...	7	1180	0	
1	7242	2.0	0	0	...	7	2170	400	
2	10000	1.0	0	0	...	6	770	0	
3	5000	1.0	0	0	...	7	1050	910	
4	8080	1.0	0	0	...	8	1680	0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
0	1955	0	98178	47.5112	-122.257	1340	
1	1951	1991	98125	47.7210	-122.319	1690	
2	1933	0	98028	47.7379	-122.233	2720	
3	1965	0	98136	47.5208	-122.393	1360	
4	1987	0	98074	47.6168	-122.045	1800	

	sqft_lot15
0	5650
1	7639
2	8062
3	5000
4	7503

[5 rows x 21 columns]

```
[55]: df_H.dtypes
```

```
[55]: id                int64
      date              object
      price             float64
      bedrooms          int64
      bathrooms         float64
      sqft_living        int64
      sqft_lot           int64
      floors             float64
      waterfront         int64
      view              int64
      condition          int64
      grade              int64
      sqft_above         int64
      sqft_basement      int64
      yr_built           int64
      yr_renovated        int64
      zipcode            int64
      lat                float64
      long               float64
      sqft_living15       int64
      sqft_lot15         int64
      dtype: object
```

```
[56]: df_H1 = df_H.copy()
      df_H1.drop("id", axis = 1, inplace = True)
      df_H1.describe()
```

```
[56]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
count	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+04	
mean	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+04	
std	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+04	
min	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02	
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	

	floors	waterfront	view	condition	grade	\
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000	
mean	1.494309	0.007542	0.234303	3.409430	7.656873	
std	0.539989	0.086517	0.766318	0.650743	1.175459	
min	1.000000	0.000000	0.000000	1.000000	1.000000	
25%	1.000000	0.000000	0.000000	3.000000	7.000000	
50%	1.500000	0.000000	0.000000	3.000000	7.000000	
75%	2.000000	0.000000	0.000000	4.000000	8.000000	

max	3.500000	1.000000	4.000000	5.000000	13.000000
-----	----------	----------	----------	----------	-----------

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1788.390691	291.509045	1971.005136	84.402258	98077.939805
std	828.090978	442.575043	29.373411	401.679240	53.505026
min	290.000000	0.000000	1900.000000	0.000000	98001.000000
25%	1190.000000	0.000000	1951.000000	0.000000	98033.000000
50%	1560.000000	0.000000	1975.000000	0.000000	98065.000000
75%	2210.000000	560.000000	1997.000000	0.000000	98118.000000
max	9410.000000	4820.000000	2015.000000	2015.000000	98199.000000

	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	47.560053	-122.213896	1986.552492	12768.455652
std	0.138564	0.140828	685.391304	27304.179631
min	47.155900	-122.519000	399.000000	651.000000
25%	47.471000	-122.328000	1490.000000	5100.000000
50%	47.571800	-122.230000	1840.000000	7620.000000
75%	47.678000	-122.125000	2360.000000	10083.000000
max	47.777600	-121.315000	6210.000000	871200.000000

```
[58]: H_floors=df_H1['floors'].value_counts()
      H_floors.to_frame()
```

```
[58]:      floors
      1.0    10680
      2.0     8241
      1.5     1910
      3.0      613
      2.5     161
      3.5       8
```

```
[59]: import matplotlib.pyplot as plt
      %matplotlib inline
      import seaborn as sns
```

```
[64]: df_H1['column_name'] = df_H1['waterfront'].astype('bool')
      df_H1.head()
```

```
[64]:      date      price  bedrooms  bathrooms  sqft_living  sqft_lot \
0  20141013T000000  221900.0         3         1.00         1180      5650
1  20141209T000000  538000.0         3         2.25         2570      7242
2  20150225T000000  180000.0         2         1.00          770     10000
3  20141209T000000  604000.0         4         3.00         1960      5000
4  20150218T000000  510000.0         3         2.00         1680      8080
```

	floors	waterfront	view	condition	...	sqft_above	sqft_basement	\
0	1.0	0	0	3	...	1180	0	
1	2.0	0	0	3	...	2170	400	
2	1.0	0	0	3	...	770	0	
3	1.0	0	0	5	...	1050	910	
4	1.0	0	0	3	...	1680	0	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
0	1955	0	98178	47.5112	-122.257	1340	
1	1951	1991	98125	47.7210	-122.319	1690	
2	1933	0	98028	47.7379	-122.233	2720	
3	1965	0	98136	47.5208	-122.393	1360	
4	1987	0	98074	47.6168	-122.045	1800	


	sqft_lot15	column_name
0	5650	False
1	7639	False
2	8062	False
3	5000	False
4	7503	False

[5 rows x 21 columns]


[77]:

```
[77]:  waterfront    price
0         0  221900.0
1         0  538000.0
2         0  180000.0
3         0  604000.0
4         0  510000.0
```

[82]: `ax = sns.boxplot(x="waterfront", y="price", data=df_H1)`



```
[83]: ax = sns.regplot(x="sqft_above", y="price", data=df_H1)
```



```
[84]: from sklearn.linear_model import LinearRegression as ln
```

```
[89]: X=df_H1[['sqft_above']]
      Y=df_H1['price']

      reg = ln().fit(X, Y)
      reg.score(X, Y)
```

```
[89]: 0.36671175283827917
```

```
[90]: X=df_H1[['floors', 'waterfront', 'lat', 'bedrooms', 'sqft_basement', 'view', 'bathrooms', 'sqft_living']]
      Y=df_H1['price']

      reg = ln().fit(X, Y)
      reg.score(X, Y)
```

```
[90]: 0.6577086983978812
```

```
[122]: from sklearn.preprocessing import PolynomialFeatures
      from sklearn.preprocessing import StandardScaler
      from sklearn.pipeline import Pipeline

      Input=[('scale',StandardScaler()), ('polynomial',PolynomialFeatures(include_bias=False)), ('model',ln())]
      pipe=Pipeline(Input)
      pipe
      pipe.fit(X,Y)
      ypipe=pipe.predict(X)
      from sklearn.metrics import r2_score
      r_squared = r2_score(Y, ypipe)
      print('The R-square value is: ', r_squared)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/preprocessing/data.py:625: DataConversionWarning: Data with
input dtype int64, float64 were all converted to float64 by StandardScaler.
```

```
    return self.partial_fit(X, y)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/base.py:465: DataConversionWarning: Data with input dtype
int64, float64 were all converted to float64 by StandardScaler.
```

```
    return self.fit(X, y, **fit_params).transform(X)
```

```
The R-square value is: 0.7513458495077826
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/pipeline.py:331: DataConversionWarning: Data with input dtype
int64, float64 were all converted to float64 by StandardScaler.
```

```
    Xt = transform.transform(Xt)
```

```
[131]: from sklearn.model_selection import train_test_split

pr=PolynomialFeatures(degree=2)

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.3,
↳random_state=0)

x_train_pr=pr.fit_transform(x_train)
x_test_pr=pr.fit_transform(x_test)
```

```
[132]: from sklearn.linear_model import Ridge
RigeModel=Ridge(alpha=0.1)
```

```
[135]: RigeModel.fit(x_train_pr, y_train)
yhat =RigeModel.predict(x_test_pr)
r_squared = r2_score(y_test, yhat)
print('The R-square value is: ', r_squared)
```

The R-square value is: 0.7307388996885178

```
[ ]:
```